

Machine Learning for Text Mining -introduction-

Antske Fokkens & Pia Sommerauer
CBS course

Supervised Machine Learning

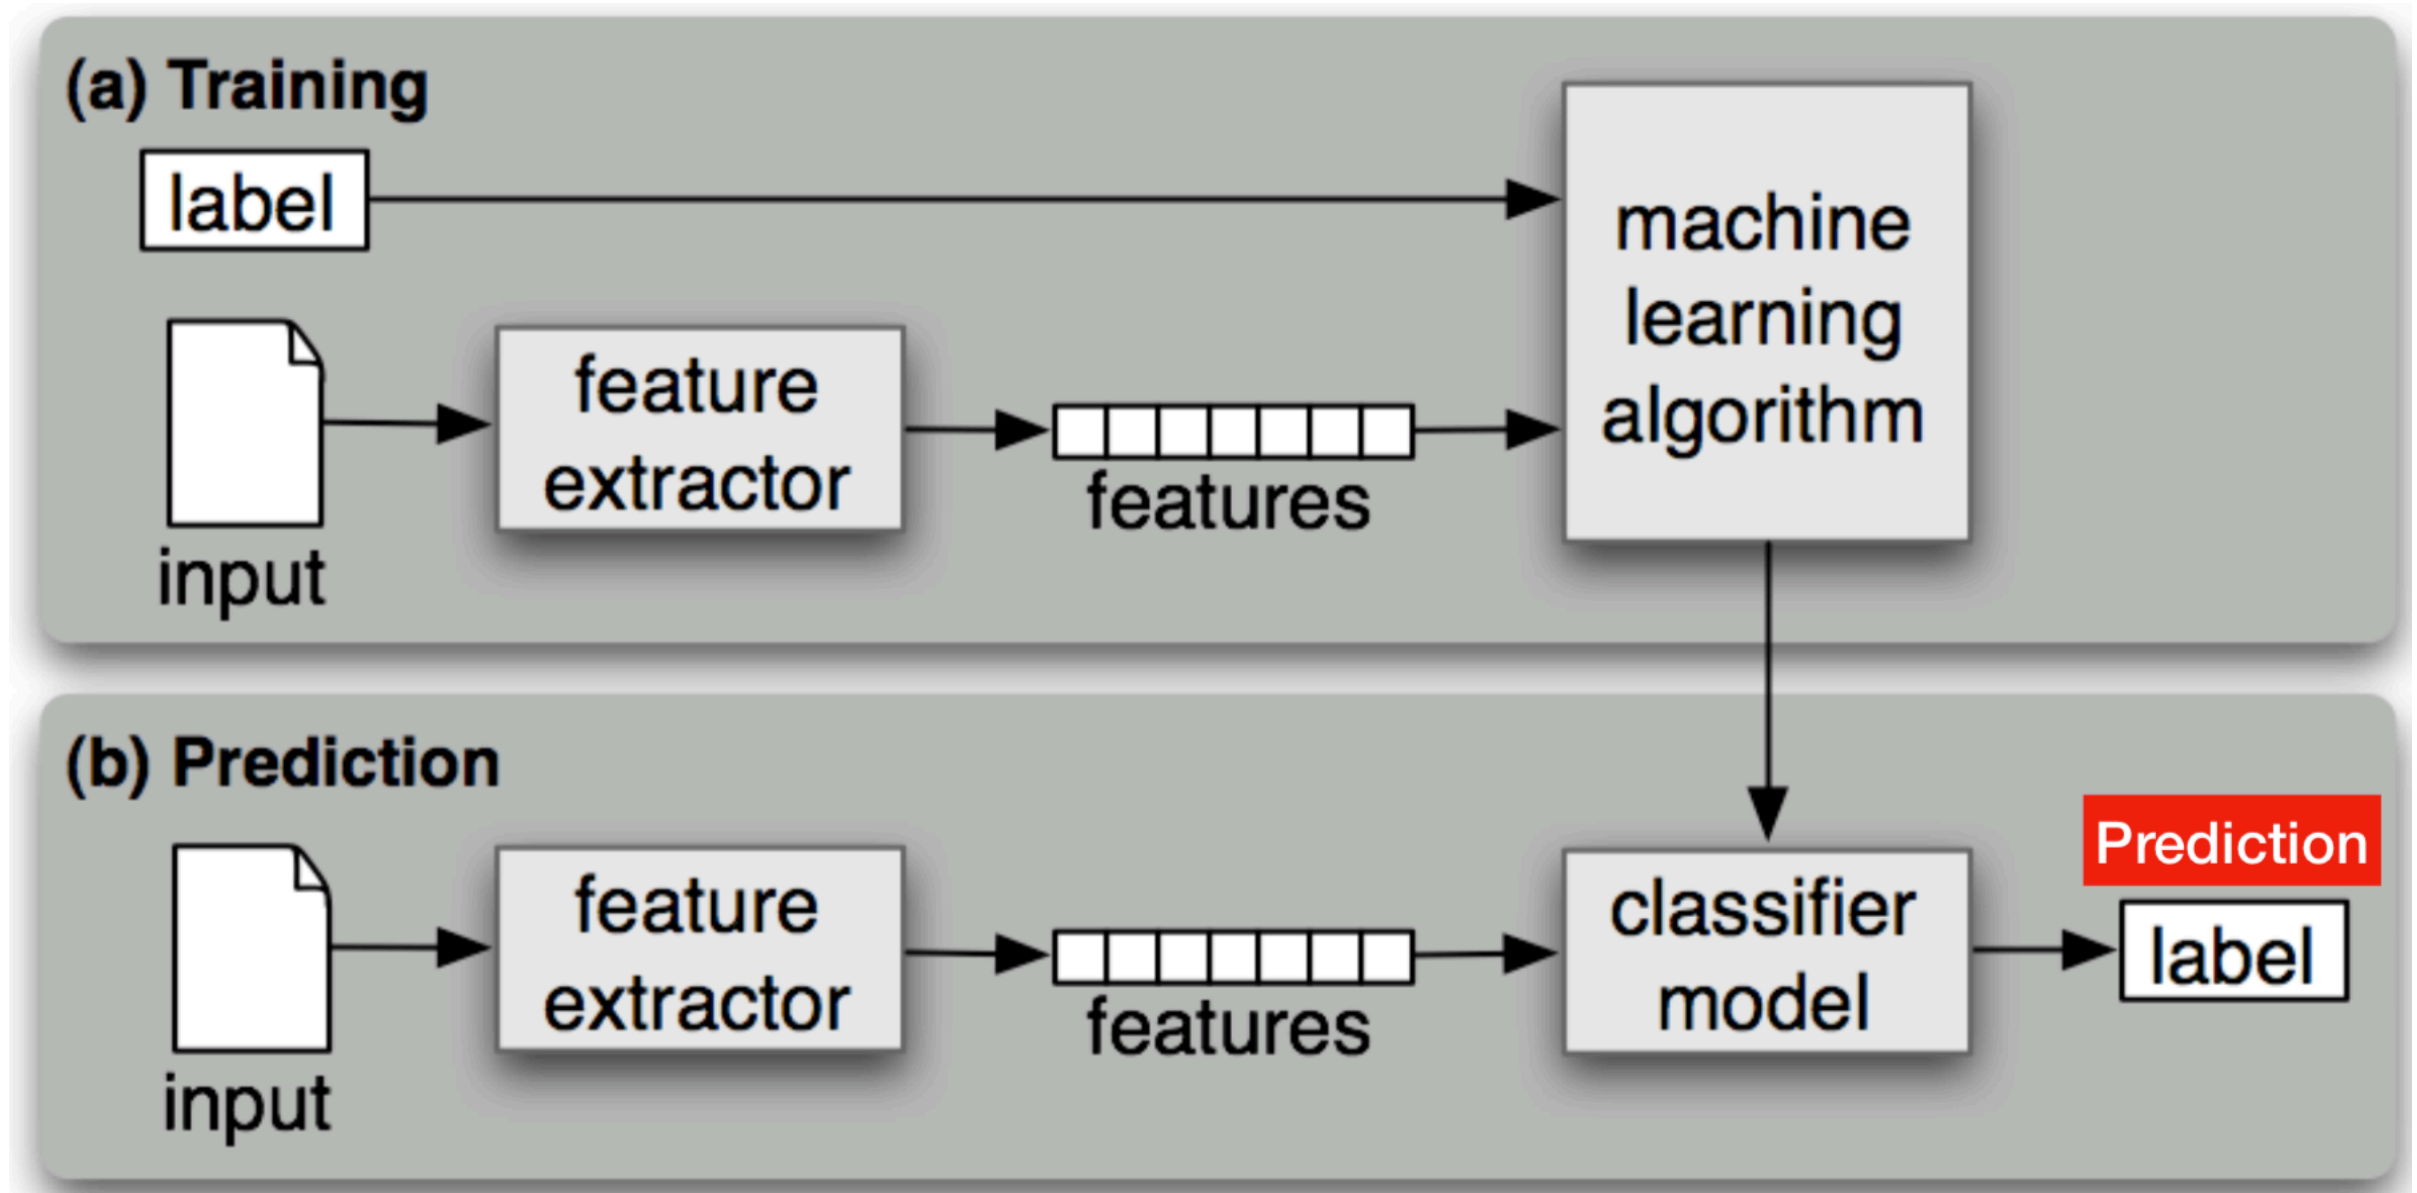



Image taken from Piek Vossen's slides (see lecture machine learning)


Translating to numbers

- Methods directly related to statistics (distribution & correlation)
- Many other forms: geometric representation & identify lines


ML methods (examples)



Source: <https://commons.wikimedia.org/wiki/File:Linear-svm-scatterplot.svg>



Source: https://en.m.wikipedia.org/wiki/File:Svm_max_sep_hyperplane_with_margin.png



Source: <https://commons.wikimedia.org/wiki/File:KnnClassification.svg>

ML: basic overview

- Many approaches:
 1. represent features as vectors with numerical values
 2. predict class based on feature vector
- For example:
 - K-nearest neighbor: pick majority class of k-nearest points in space
 - Logistic regression: find hyperplane that separates the data best (minimizing loss)


Features

- Common features (for many tasks):
 - POS-tag
 - Word
 - Lemma
 - ngrams
- More advanced:
 - Chunks
 - Syntactic dependencies
 - Word sense

Features

- Common features (for many tasks):

- POS-tag



- More advanced:


- Chunks
 - Syntactic dependencies
 - Word sense

Input representation

- Many Machine Learning Classifiers use **input representations** that correspond to vectors
- How can we translate linguistic information to space?

Representations

- A number in one dimension, e.g. size of houses: represent size as number (e.g. square meters)
- Different dimensions allocated to various values, e.g. blood type of a person: represent value in a 4-dimensional vector




Linguistic features

How to represent?

- POS-tags
- Words/lemmas
- Chunks
- Dependencies


Penn Treebank Pos-tags

- 36 tags: 36 dimensions (one-hot representation)

[illegible]

Word - Lemmas

- One-hot representations: vocabulary-size dimensions

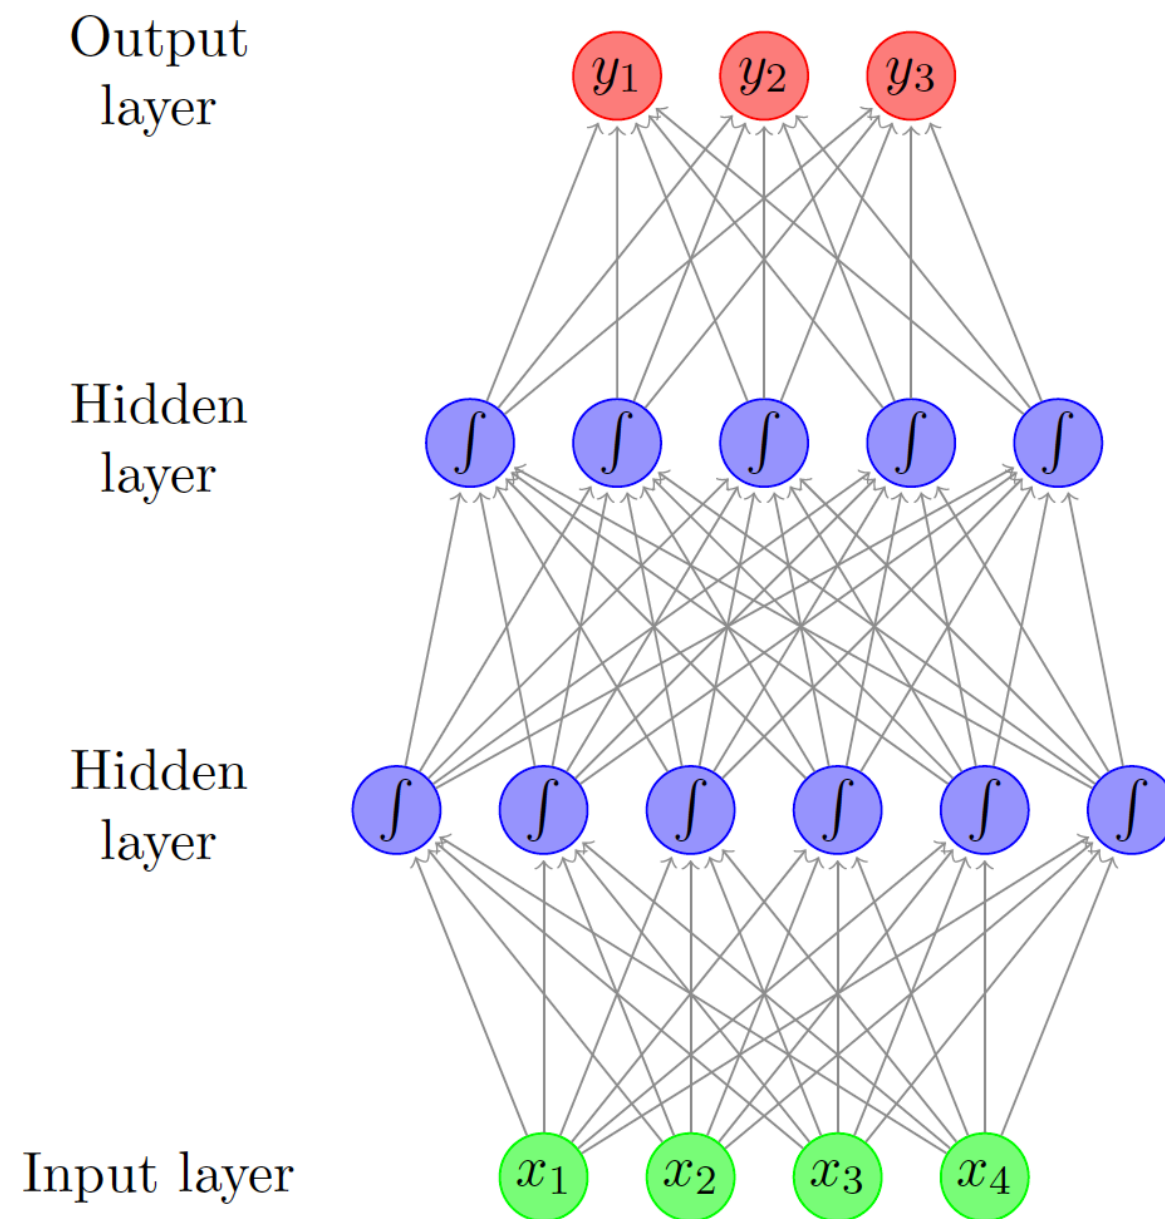


Word - Lemmas

- Word embeddings
 - Lower-dimensional (~50-1,000), higher density representations
 - Presentation based on context a word occurs in
- => words in similar context have similar representation
- => algorithm can gather evidence from similar words in training data

MORE ON WORD EMBEDDINGS LATER !

A Basic Neural network



from Goldberg (2015)

Brain Analogy





Figure 3: Illustration of a biological neuron (left) and its mathematical model (right) [2].

+ non-linear function

By: Andrej Karapthy 2015

A Basic Neural network



from Goldberg (2015)

Requirements

- workable size of feature space:
 - limited feature space
 - => new representations of words
 - => using research on lexical semantics

Word Embeddings

PPMI

- $\text{PPMI} = \text{argmax}(0, \log(\frac{P(w1, w2)}{P(w1)P(w2)}))$

PPMI

- High dimensional (size of vocabulary)
- Low density (zeros for all context-words occurring less than by chance)
- Relatively high impact of low frequency words

Singular Value Decomposition (SVD)

- Method to reduce the number of dimensions:

given a $m \times n$ matrix, construct a $m \times k$ matrix,
where $k \ll n$

- Uses linear algebra to reduce the number of dimensions,
preserving most of the variance of the original matrix

SVD

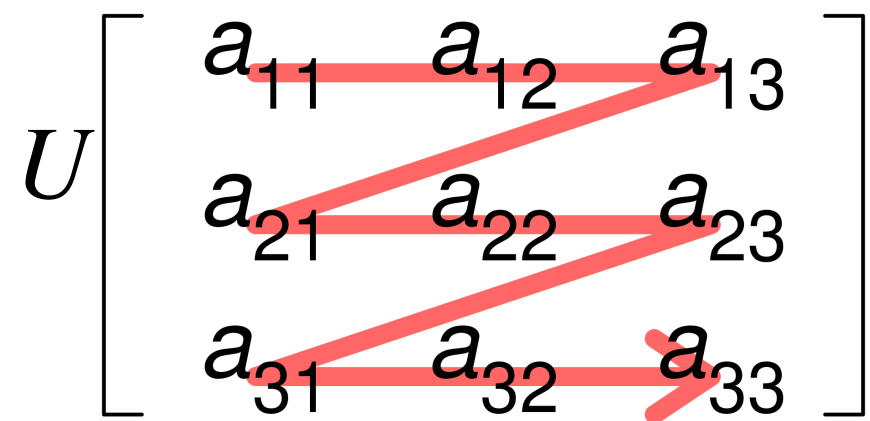
- A matrix A can be broken down (decomposed) into the product of three matrices:

$$A = U\Sigma V^T$$

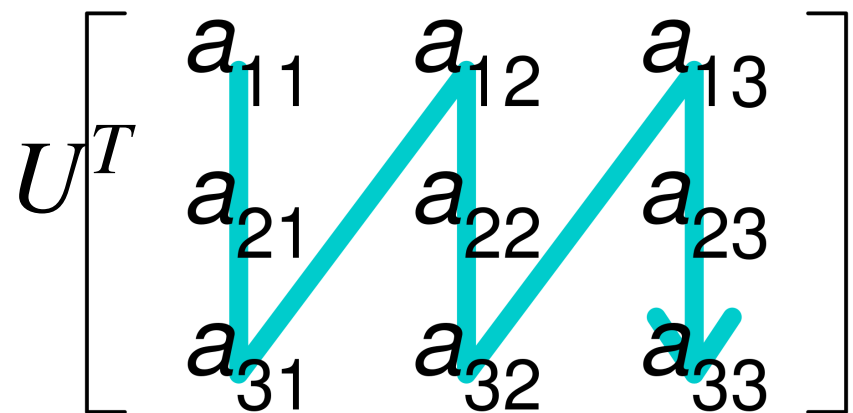
- where U and V are orthogonal
- The columns of U are orthonormal eigenvectors of AA^T
- The columns of V are orthonormal eigenvectors of $A^T A$
- Σ is a diagonal matrix containing square roots of eigenvalues from U or V in descending order

U and V are orthogonal

Row-major order

$$U \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$
A diagram of a 3x3 matrix U. Red arrows indicate a row-major traversal pattern: starting at a11, moving right to a12 and a13, then down to a21, a22, a23, then down to a31, a32, and finally a33. The arrow from a32 to a33 is crossed out with a red 'X'.


Column-major order

$$U^T \begin{bmatrix} a_{11} & a_{12} & a_{13} \\ a_{21} & a_{22} & a_{23} \\ a_{31} & a_{32} & a_{33} \end{bmatrix}$$
A diagram of a 3x3 matrix U^T. Cyan arrows indicate a column-major traversal pattern: starting at a11, moving down to a21 and a31, then up-right to a12, a22, a32, then up-right to a13, a23, and finally a33.

$$UU^T = I$$


$$\begin{bmatrix} 1 & 0 & 0 \\ 0 & 1 & 0 \\ 0 & 0 & 1 \end{bmatrix}$$

Capturing most variance




from Baroni & Boleda

Capturing most variance




from Baroni & Boleda

Capturing most variance




from Baroni & Boleda

Capturing most variance




from Baroni & Boleda

Capturing most variance




from Baroni & Boleda

Capturing most variance



from Baroni & Boleda

Capturing most variance



from Baroni & Boleda

Reducing Dimensions

- Columns of U and V are ordered according to highest associated eigenvalue
- Diagonal values of Σ are ordered starting with highest (root of) eigenvalues

=> Using the first d rows of U , the first d columns of V^T and $d \times d$ rows and columns of Σ guarantees that we end up with those values that provide the highest variance.

Predictive Models

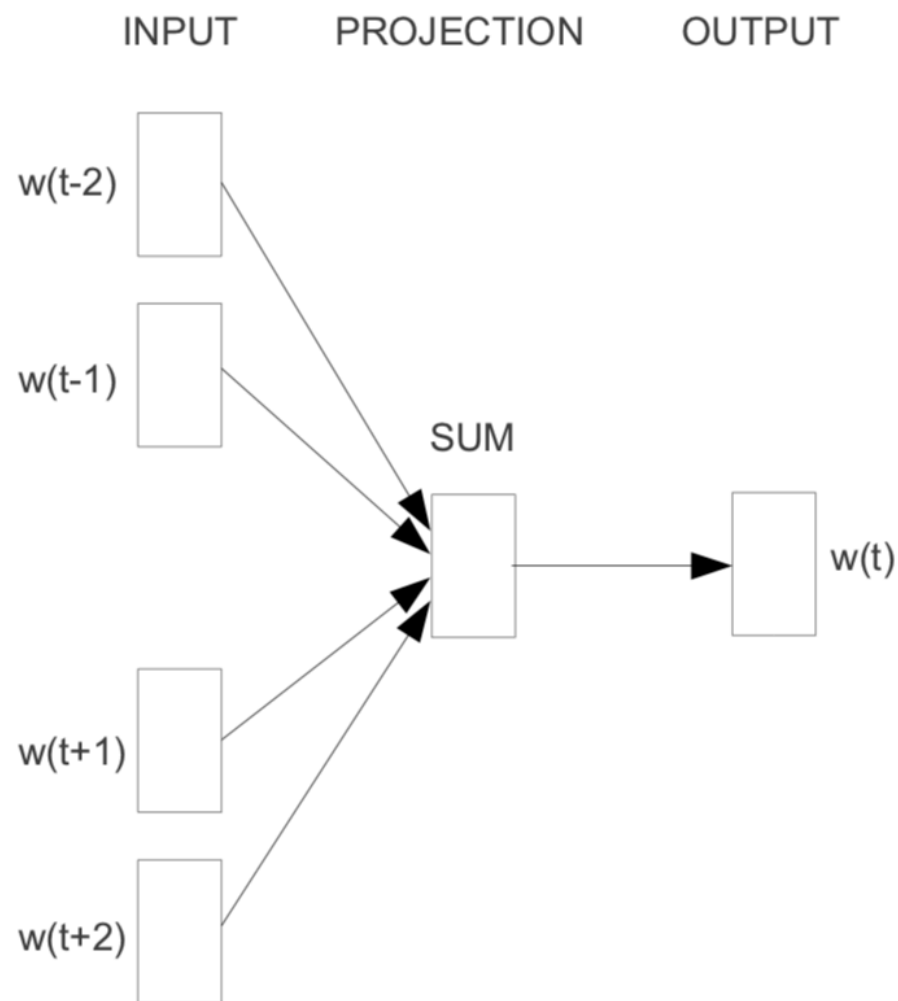
- PPMI: captures desirable properties, but is still sparse and high dimensional
- SVD: generalizes better is high density and lower dimensions, but is inefficient to obtain (and does not perform perfectly)
- Alternative idea: use language modeling as an auxiliary task for creating word embeddings

=> machine learning to **predict** which words occur in each other's context


Predictive Models

- word2vec (Mikolov et al. 2013a,b):
 - Several methods for creating embeddings
 - All start from randomly initiated vectors with a preset number of dimensions
 - Two vocabulary matrices: one for the words, one for contexts
 - Models: CBOW & Skipgram
 - Training: hierarchical softmax & negative sampling
 - Preprocessing: dynamic context windows, subsampling, delete rare words

CBOW vs Skipgram



CBOW



Skip-gram

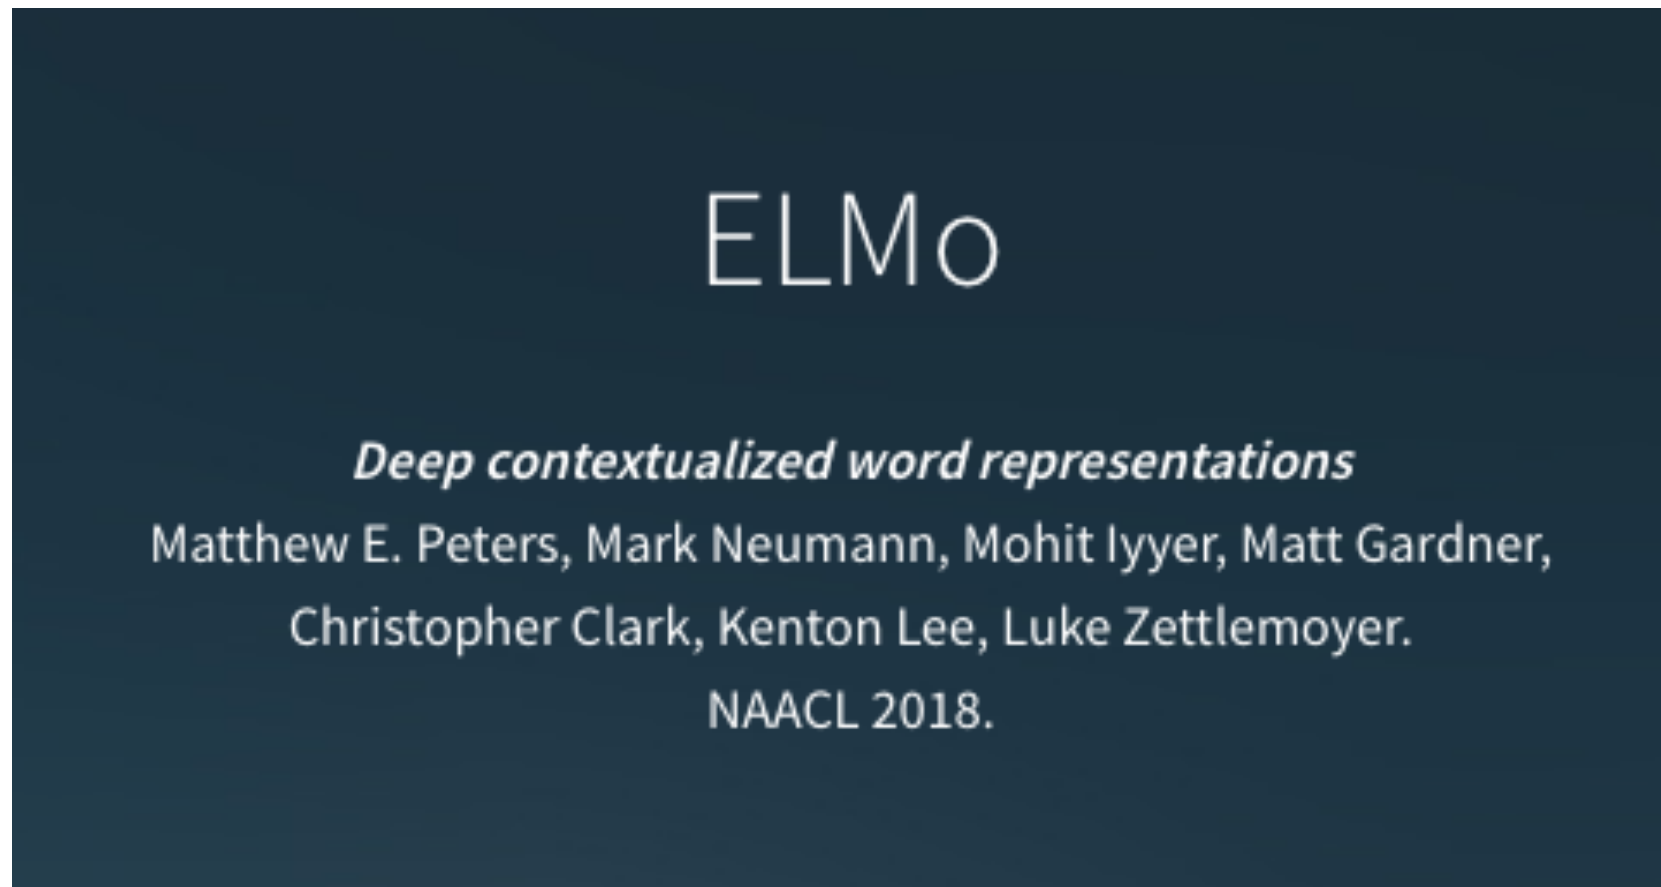
Training

- Hierarchical softmax: efficient way to determine most probable context given a word (or vice versa) over the whole model
- Negative sampling: distinguish the actual context words from k other words (randomly chosen)

Result

- d -dimensional word & context embeddings
- distributional model: the vectors representing words are kept

Latest Developments



<https://allennlp.org/elmo>

- Easy to use (with advanced context selection & neural network for learning)
- State-of-the-art for many NLP tasks

Word Embeddings as features

- Suitable to be used as input in neural networks
- Generalize better => enhance many tasks, also when used with other machine learning approaches
- Ongoing research: which embeddings (how created) for which task?

=> dependency parsing

=> sentiment analysis

Sources

- Baroni & Boleda. Distributional Semantic Models <https://www.cs.utexas.edu/~mooney/cs388/slides/dist-sem-intro-NLP-class-UT.pdf>
- Goldberg, Yoav (2015) [A primer on neural network models for Natural Language Processing](#)
- Mikolov, T., K. Chen, G. Corrado and J. Dean (2013) Efficient estimation of word representations in vector space. <https://arxiv.org/pdf/1301.3781.pdf>
- Shaffy, Athif (2017) Vector Representation of Text for Machine Learning. <https://medium.com/@athif.shaffy/one-hot-encoding-of-text-b69124bef0a7>
- Pia Sommerauer. What is in a word embedding vector?
- All reported sources were accessed between January 14-16 2019